

HELLO NODE KUBERNETES:

Create the backend and run using node servername

Create Docker Container Image:

- .. write a docker file
- .. build the image
- .. run the container and see the backend running.
- .. push the working image to docker hub/ in this lab
- .. its google artifact registry..

This thing need attention? `gcloud artifacts repositories create my-docker-repo \ --repository-format=docker \ --location=us-west1 \ --project=qwiklabs-gcp-02-404ec7e1e03f..`

How can this be used locally?

Then this..

`gcloud auth configure-docker`

Then ..

`docker tag hello-node:v1`

`us-west1-docker.pkg.dev/project_id/my-docker-repo/hello-node:v1`

Followed by 😊

`docker push`

`us-west1-docker.pkg.dev/project_id/my-docker-repo/hello-node:v1`

Now you have a project-wide Docker image available which Kubernetes can access and orchestrate.

1. Create Your Kubernetes Cluster

A Kubernetes cluster provides the environment to run your containerized applications. It consists of a **control plane** (the "master") managed by Google and **worker nodes**, which are the virtual machines that actually run your application containers.

First, you need to tell the `gcloud` command-line tool which Google Cloud project you're working with.

Bash

```
# Set your active project
```

```
gcloud config set project YOUR_PROJECT_ID
```

Next, create the cluster. This command provisions the necessary compute resources.

Bash

```
# Create a GKE cluster named 'hello-world'  
gcloud container clusters create hello-world \  
  --num-nodes 2 \  
  --machine-type e2-medium \  
  --zone "us-west1-b"
```

- **--num-nodes 2**: Creates a cluster with **two worker nodes**.
- **--machine-type e2-medium**: Specifies the size (CPU/RAM) of each worker node.
- **--zone "us-west1-b"**: Defines the physical location where your cluster will be created.

2. Create Your Pods via a Deployment

You don't usually create Pods directly. Instead, you create a **Deployment**, which is a Kubernetes object that manages a set of identical Pods. It ensures that a specified number of Pods are always running and handles updates and rollbacks.

Bash

```
# Create a deployment named 'hello-node' from a container image  
kubectl create deployment hello-node \  
  --image=us-west1-docker.pkg.dev/YOUR_PROJECT_ID/my-docker-repo/hello-no  
  de:v1
```

- This command tells Kubernetes to pull the specified container image and run it in a Pod managed by the **hello-node** deployment.

3. Inspect and Troubleshoot Your Application

Once you've created resources, you need to verify they are running correctly.

Key Inspection Commands:

Bash

```
# Check the status of your deployments
```

```
kubectl get deployments
```

```
# See the individual pods created and managed by the deployment
```

```
kubectl get pods
```

```
# Get a high-level overview of your cluster's health and endpoints
```

```
kubectl cluster-info
```

```
# View your current kubectl configuration (e.g., which cluster you're connected to)
```

```
kubectl config view
```

Key Troubleshooting Commands:

Bash

```
# See all events in the cluster (useful if a pod is stuck or failing)
```

```
kubectl get events
```

```
# View the logs from a specific pod to debug application errors
```

```
kubectl logs <pod-name>
```

4. Allow External Traffic

By default, your application is only accessible from within the cluster. To expose it to the internet, you create a **Service**. A Service provides a stable network endpoint (IP address and port) for your set of Pods.

Bash

```
# Expose the 'hello-node' deployment using a LoadBalancer Service  
kubectl expose deployment hello-node --type="LoadBalancer" --port=8080
```

- `--type="LoadBalancer"`: This is crucial. It tells the cloud provider (GCP) to provision an external network load balancer, which will get a public IP address and route external traffic to your Pods.

To find the public IP address, run:

Bash

```
# Get information about your services, including the external IP  
kubectl get services
```

Look for the **EXTERNAL-IP** address. It might take a few minutes to appear. Once it does, you can access your app at `http://<EXTERNAL_IP>:8080`.

5. Scale Up Your Service

To handle more traffic or ensure high availability, you can easily increase the number of Pods running your application.

Bash

```
# Scale the number of replicas in the deployment to 4
```

```
kubectl scale deployment hello-node --replicas=4
```

```
# Verify that 4 pods are now running
```

```
kubectl get pods
```

The Deployment and the LoadBalancer Service will automatically handle distributing traffic across all four pods.

6. Roll Out an Upgrade

To update your application to a new version, you first need to build and push your new container image ([v2](#)) to the registry.

Bash

1. Tag the new version of your image

```
docker tag hello-node:v2  
us-west1-docker.pkg.dev/YOUR_PROJECT_ID/my-docker-repo/hello-node:v2
```

2. Push the new image to the Artifact Registry

```
docker push  
us-west1-docker.pkg.dev/YOUR_PROJECT_ID/my-docker-repo/hello-node:v2
```

Now, you can tell Kubernetes to perform a **rolling update**. This updates the Pods one by one without any downtime for your users.

Bash

3. Update the deployment to use the new image version

```
kubectl set image deployment/hello-node  
hello-node=us-west1-docker.pkg.dev/YOUR_PROJECT_ID/my-docker-repo/hello  
-node:v2
```

Kubernetes will terminate old Pods and create new ones with the [v2](#) image until the entire deployment is updated.



Local Solution: Running This on Your Linux Terminal

You asked what can be done without `gcloud`. You can replicate this entire workflow on your local machine! The `kubectl` commands are universal, but you need a local Kubernetes cluster instead of GKE.

Popular tools for local Kubernetes clusters include **Minikube**, **Kind**, and **Docker Desktop's built-in Kubernetes**. Here's the workflow using **Minikube**.

Step 1: Install a Local Cluster

First, install `kubectl`, Docker, and Minikube. Then, start your local cluster:

Bash

```
# Start a single-node Kubernetes cluster on your machine  
minikube start
```

Step 2: Build & Load Your Docker Image

Instead of pushing your image to a remote registry like GCR, you can build it locally and load it directly into your Minikube cluster's Docker environment.

Bash

```
# This command points your terminal's Docker daemon to the one inside  
Minikube  
  
eval $(minikube docker-env)
```

```
# Now, build your image. It will be immediately available to the cluster.
```

```
docker build -t hello-node:v1 .
```

```
# To switch back to your host's Docker daemon when you're done:
```

```
eval $(minikube docker-env -u)
```

Step 3: Run `kubectl` Commands (They're the Same!)

All the `kubectl` commands you used before will work exactly the same way against your local cluster.

Bash

```
# Create the deployment using the locally built image  
kubectl create deployment hello-node --image=hello-node:v1
```

```
# Inspect it
```

```
kubectl get pods
```

```
# Scale it
```

```
kubectl scale deployment hello-node --replicas=4
```

Step 4: Expose The Service (The Main Difference)

A local environment doesn't have a cloud load balancer. So, `type="LoadBalancer"` behaves differently. The two common ways to access your service locally are:

Using `minikube service` (Easiest): This command automatically creates a tunnel and opens the service URL in your browser. It's the recommended way for Minikube.

Bash

```
# Expose the service. Minikube will handle the LoadBalancer type.
```

```
kubectl expose deployment hello-node --type="LoadBalancer" --port=8080
```

```
# Let Minikube open the service for you
```

```
minikube service hello-node
```

1.

Using NodePort: This exposes the service on a specific port on the cluster's node IP.

Bash

```
# Expose using a NodePort instead of LoadBalancer
```

```
kubectl expose deployment hello-node --type=NodePort --port=8080
```

```
# Find the URL for the service
```

```
minikube service hello-node --url
```

2.

This local setup is perfect for development, testing, and learning Kubernetes fundamentals without incurring any cloud costs.