



Understanding Google Cloud Monitoring (with Local Parallels)

This guide is designed to **translate a Google Cloud Monitoring Lab** into something meaningful — helping you understand:

-  What each task in the lab *really does*
 -  Why it matters in real-world DevOps
 -  How to replicate it **locally** using free, open-source tools like Docker, Prometheus, Grafana, and Blackbox Exporter
-



Task 1: Creating Virtual Machines (VMs)

What you're doing:

Creating a Compute Engine instance (`instance2`) inside Project 2 — essentially a Linux server hosted by Google Cloud.

Compute engine: a virtual machine or bare metal server hosted on Google's infrastructure, providing on demand computing power to run apps and workload.

Why it matters:

Monitoring starts with something to observe. Having multiple instances lets you simulate different environments like `staging`, `testing`, and `production`.

Do this locally:

Using **Docker**, you can spin up two Linux-like containers:

```
docker run -d --name instance1 ubuntu sleep infinity
```

```
docker run -d --name instance2 ubuntu sleep infinity
```

`sleep infinity`: This is the command that runs inside the container once it starts.

- `sleep` is a Unix command that pauses for a specified amount of time.
- `infinity` is an argument telling it to sleep forever.
- This is a common technique used to keep a basic container running indefinitely. By default, a Docker container stops as soon as its main command finishes. Since `sleep infinity` never finishes, the container stays alive, allowing you to `docker exec` into it to run commands or use it for other purposes.

Both containers now represent your “VMs.”

Task 2: Setting Up a Monitoring Metrics Scope

What you're doing:

Creating a **Metrics Scope** that unifies metrics from multiple projects.

Why it matters:

Companies often have different GCP projects for different environments. Metrics Scopes give a *single dashboard* to monitor all.

Do this locally (Prometheus Federation):

```
scrape_configs:  
  - job_name: 'project1'  
  
    static_configs:  
      - targets: ['localhost:9100']  
  
  - job_name: 'project2'  
  
    static_configs:  
      - targets: ['localhost:9200']
```

This is like combining metrics from multiple Prometheus servers into one view.

Task 3: Creating a Cloud Monitoring Group

What you're doing:

Defining a group ([DemoGroup](#)) that includes all VMs whose names contain “instance.”

Why it matters:

Groups organize related resources for focused alerting or visualization — like grouping all frontend nodes.

Do this locally:

In [Grafana](#), you can create a dashboard that filters instances using queries like:

```
up{job=~"instance.*"}
```

Task 4: Uptime Checks

What you're doing:

Configuring an automated **Uptime Check** to see if your instance’s port 22 (SSH) is open.

Why it matters:

It’s a fundamental “is this alive?” check for all services — helps detect global outages.

Do this locally (Blackbox Exporter):

```
modules:  
  tcp_22:  
    prober: tcp  
    tcp:  
      query_response:
```

Prometheus scrapes this exporter to verify if a target port or endpoint is reachable.



Task 5: Alerting Policies

What you're doing:

Creating an alert that triggers when the uptime check fails.

Why it matters:

Automation in alerting ensures downtime doesn't go unnoticed — critical for reliability.

Do this locally (Prometheus Rule):

groups:

```
- name: example
```

rules:

```
- alert: InstanceDown
```

```
expr: up == 0
```

```
for: 1m
```

labels:

```
severity: critical
```

annotations:

```
summary: "Instance {{ $labels.instance }} is down"
```

How to see it:

When Prometheus detects a container stopped, this alert fires.

If connected to Alertmanager, it can notify via email, Slack, or Discord.



Task 6: Custom Dashboards

What you're doing:

Building a dashboard in Cloud Monitoring to visualize uptime metrics.

Why it matters:

Dashboards help teams *see patterns and diagnose issues visually* — crucial for production observability.

Do this locally (Grafana):

1. Open Grafana (<http://localhost:3000>)
2. Add Prometheus as a data source.
3. Create a new dashboard → “Add Panel.”

Use query:

`up`

- 4.
 5. Select Line or Gauge visualization.
-



Task 7: Causing an Incident (Simulating a Failure)

What you’re doing:

Stopping your instance to intentionally fail an uptime check — generating an alert/incident.

Why it matters:

Testing failure conditions validates your monitoring system.

It’s an early lesson in **chaos engineering** — breaking things to improve resilience.

Do this locally:

`docker stop instance2`

Prometheus marks it `up == 0`.

Restart to recover:

`docker start instance2`

Summary Table

Concept	Google Cloud Component	Local / Open Source Equivalent
Virtual Machines	Compute Engine	Docker / VirtualBox
Metrics Collection	Cloud Monitoring	Prometheus
Visualization	Dashboards	Grafana
Alerting	Alert Policy	Prometheus + Alertmanager
Uptime Check	Cloud Uptime	Blackbox Exporter
Incident Handling	Cloud Incidents	Alertmanager UI

Setting Up a Local Monitoring Stack (Full Hands-on)

Here's how to replicate **Google Cloud Monitoring** locally in one go, using Docker.

Step 1: Create a `docker-compose.yml` file

```
version: '3'
```

```
services:
```

```
  prometheus:
```

```
    image: prom/prometheus
```

```
    container_name: prometheus
```

```
volumes:
  - ./prometheus.yml:/etc/prometheus/prometheus.yml

ports:
  - "9090:9090"

grafana:
  image: grafana/grafana
  container_name: grafana
  ports:
    - "3000:3000"
  environment:
    - GF_SECURITY_ADMIN_USER=admin
    - GF_SECURITY_ADMIN_PASSWORD=admin

blackbox:
  image: prom/blackbox-exporter
  container_name: blackbox
  ports:
    - "9115:9115"
  volumes:
    - ./blackbox.yml:/etc/blackbox_exporter/config.yml
```

Step 2: Create the Prometheus Config (`prometheus.yml`)

```
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'prometheus'
    static_configs:
      - targets: ['prometheus:9090']

  - job_name: 'blackbox'
    metrics_path: /probe
    params:
      module: [tcp_22]
    static_configs:
      - targets:
          - instance1
          - instance2

  relabel_configs:
    - source_labels: [__address__]
      target_label: __param_target
    - target_label: instance
      replacement: blackbox
    - target_label: __address__
      replacement: blackbox:9115
```

Step 3: Create the Blackbox Config (`blackbox.yml`)

`modules:`

```
  tcp_22:  
    prober: tcp  
    timeout: 5s
```

Step 4: Launch Everything

Run:

```
docker-compose up -d
```

Check:

- Prometheus → <http://localhost:9090>
 - Grafana → <http://localhost:3000>
 - Blackbox → <http://localhost:9115>
-

Step 5: Add Grafana Dashboard

1. Log in: `admin / admin`
2. Add data source → choose “Prometheus” → URL: <http://prometheus:9090>
3. Create Dashboard → Add Panel

Use queries like:

```
up  
probe_success
```

4. These are your uptime metrics!
-

Step 6: Simulate a Failure

Stop one of your test containers:

```
docker stop instance2
```

Watch in:

- **Prometheus:** `up` value turns 0
- **Grafana:** dashboard shows drop
- **Alerts:** trigger via Prometheus rule if configured

Restart:

```
docker start instance2
```

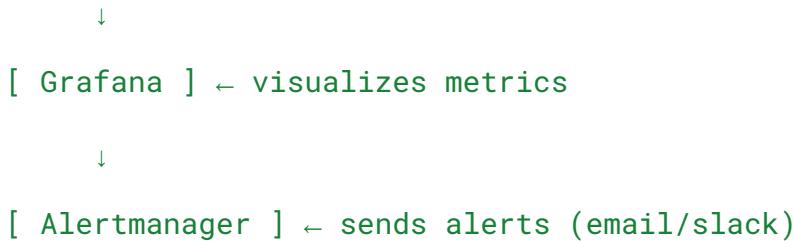
Within a minute, it clears — just like a closed incident in GCP.

Local Architecture Overview

[Docker Containers]

↓ metrics

[Prometheus] ← scrapes uptime



This replicates:

- GCP's Compute Engine → Docker containers
- Cloud Monitoring → Prometheus
- Cloud Dashboards → Grafana
- Uptime Checks → Blackbox Exporter
- Alert Policies → Prometheus alert rules
- Incidents → Alertmanager notifications

🎯 Final Takeaway

What you learned here mirrors **real-world SRE (Site Reliability Engineering)** concepts:

- Always **instrument** systems to emit metrics
- Build **visibility** using dashboards
- Automate **alerting** for failures
- Validate monitoring by **causing test incidents**

Cloud or local — the fundamentals are identical.

Once you grasp this, tools like GCP, AWS CloudWatch, or Azure Monitor all follow the same blueprint.